


```
import numpy as np
import pandas as pd
```

Importing Data

```
crop = pd.read_csv("Crop_recommendation.csv")
crop.head()
```



	N	P	K	temperature	humidity	ph	rainfall	label
0	90	42	43	20.879744	82.002744	6.502985	202.935536	rice
1	85	58	41	21.770462	80.319644	7.038096	226.655537	rice
2	60	55	44	23.004459	82.320763	7.840207	263.964248	rice
3	74	35	40	26.491096	80.158363	6.980401	242.864034	rice
4	78	42	42	20.130175	81.604873	7.628473	262.717340	rice

Start coding or [generate](#) with AI.

```
crop.shape
```

(2200, 8)

```
crop.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2200 entries, 0 to 2199
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   N                2200 non-null   int64
1   P                2200 non-null   int64
2   K                2200 non-null   int64
3   temperature      2200 non-null   float64
4   humidity         2200 non-null   float64
5   ph               2200 non-null   float64
6   rainfall         2200 non-null   float64
7   label            2200 non-null   object
dtypes: float64(4), int64(3), object(1)
memory usage: 137.6+ KB
```

```
crop.isnull().sum()
```

```
N          0
P          0
K          0
temperature 0
humidity    0
ph          0
rainfall    0
label       0
dtype: int64
```

```
crop.duplicated().sum()
```

0

```
crop.describe()
```

	N	P	K	temperature	humidity	ph	ra
count	2200.000000	2200.000000	2200.000000	2200.000000	2200.000000	2200.000000	2200.000000
mean	50.551818	53.362727	48.149091	25.616244	71.481779	6.469480	103.000000
std	36.917334	32.985883	50.647931	5.063749	22.263812	0.773938	54.000000
min	0.000000	5.000000	5.000000	8.825675	14.258040	3.504752	20.000000
25%	21.000000	28.000000	20.000000	22.769375	60.261953	5.971693	64.000000
50%	37.000000	51.000000	32.000000	25.598693	80.473146	6.425045	94.000000
75%	84.250000	68.000000	49.000000	28.561654	89.948771	6.923643	124.000000
max	140.000000	145.000000	205.000000	43.675493	99.981876	9.935091	298.000000


Exploring Data

```
corr = crop.corr()
corr

<ipython-input-16-1dcd70030872>:1: FutureWarning: The default value of numeric_only in I
corr = crop.corr()
```

	N	P	K	temperature	humidity	ph	rainfall
N	1.000000	-0.231460	-0.140512	0.026504	0.190688	0.096683	0.059020
P	-0.231460	1.000000	0.736232	-0.127541	-0.118734	-0.138019	-0.063839
K	-0.140512	0.736232	1.000000	-0.160387	0.190859	-0.169503	-0.053461
temperature	0.026504	-0.127541	-0.160387	1.000000	0.205320	-0.017795	-0.030084
humidity	0.190688	-0.118734	0.190859	0.205320	1.000000	-0.008483	0.094423
ph	0.096683	-0.138019	-0.169503	-0.017795	-0.008483	1.000000	-0.109069
rainfall	0.059020	-0.063839	-0.053461	-0.030084	0.094423	-0.109069	1.000000

```
import seaborn as sns
sns.heatmap(corr,annot=True,cbar=True,cmap='coolwarm')
```



```
crop['label'].value_counts()
```

```

rice      100
maize     100
jute      100
cotton    100
coconut   100
papaya    100
orange    100
apple     100
muskmelon 100
watermelon 100
grapes    100
mango     100
banana    100
pomegranate 100
lentil    100
blackgram 100
mungbean  100
mothbeans 100
pigeonpeas 100
kidneybeans 100
chickpea  100
coffee   100
Name: label, dtype: int64

```

```
crop['label'].isnull().sum()
```

```
0
```

```

import matplotlib.pyplot as plt
sns.distplot(crop["N"])
plt.show()

```

```
<ipython-input-93-c29d4605fd08>:2: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(crop["N"])
```



Encoding

```
crop_dict = {
    'rice': 1,
    'maize': 2,
    'jute': 3,
    'cotton': 4,
    'coconut': 5,
    'papaya': 6,
    'orange': 7,
    'apple': 8,
    'muskmelon': 9,
    'watermelon': 10,
    'grapes': 11,
    'mango': 12,
    'banana': 13,
    'pomegranate': 14,
    'lentil': 15,
    'blackgram': 16,
    'mungbean': 17,
    'mothbeans': 18,
    'pigeonpeas': 19,
    'kidneybeans': 20,
    'chickpea': 21,
    'coffee': 22
}
crop['crop_num'] = crop['label'].map(crop_dict)
```

```
crop['crop_num'].value_counts()
```

```
1    100
2    100
3    100
4    100
5    100
6    100
7    100
8    100
9    100
10   100
11   100
12   100
13   100
14   100
15   100
16   100
17   100
18   100
19   100
20   100
21   100
22   100
Name: crop_num, dtype: int64
```

```
#crop.drop('label',axis=1,inplace=True)
crop.head()
```

	N	P	K	temperature	humidity	ph	rainfall	label	crop_num
0	90	42	43	20.879744	82.002744	6.502985	202.935536	rice	1
1	85	58	41	21.770462	80.319644	7.038096	226.655537	rice	1
2	60	55	44	23.004459	82.320763	7.840207	263.964248	rice	1
3	74	35	40	26.491096	80.158363	6.980401	242.864034	rice	1
4	78	42	42	20.130175	81.604873	7.628473	262.717340	rice	1

Train Test Split

```
X= crop.drop(['crop_num', 'label'],axis=1)
y =crop['crop_num']
```

```
X.shape
```

```
(2200, 7)
```

```
y.shape
```

```
(2200,)
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2,random_state=42)
```

```
X_train.shape
```

```
(1760, 7)
```

```
X_test.shape
```

```
(440, 7)
```

Scale the features using MinMaxScaler

```
from sklearn.preprocessing import MinMaxScaler
ms = MinMaxScaler()
```

```
ms.fit(X_train)
X_train= ms.transform(X_train)
X_test = ms.transform(X_test)
```

```
X_train
```

```
array([[0.12142857, 0.07857143, 0.045      , ..., 0.9089898 , 0.48532225,
        0.29685161],
       [0.26428571, 0.52857143, 0.07      , ..., 0.64257946, 0.56594073,
        0.17630752],
       [0.05      , 0.48571429, 0.1      , ..., 0.57005802, 0.58835229,
        0.08931844],
       ...,
       [0.07857143, 0.22142857, 0.13      , ..., 0.43760347, 0.46198144,
        0.28719815],
       [0.07857143, 0.85      , 0.995     , ..., 0.76763665, 0.44420505,
        0.18346657],
       [0.22857143, 0.52142857, 0.085     , ..., 0.56099735, 0.54465022,
        0.11879596]])
```

```
y_train.isnull().sum()
```

```
0
```

Standarization

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
```

```
sc.fit(X_train)
X_train= sc.transform(X_train)
X_test = sc.transform(X_test)
```

```
X_train
```

```
array([[ -9.03426596e-01, -1.12616170e+00, -6.68506601e-01, ...,
         9.36586183e-01,  1.93473784e-01,  5.14970176e-03],
       [-3.67051340e-01,  7.70358846e-01, -5.70589522e-01, ...,
        -1.00470485e-01,  8.63917548e-01, -6.05290566e-01],
```

```
[ -1.17161422e+00,  5.89737842e-01, -4.53089028e-01, ...,
  -3.82774991e-01,  1.05029771e+00, -1.04580687e+00],
...,
[ -1.06433917e+00, -5.24091685e-01, -3.35588533e-01, ...,
  -8.98381379e-01, -6.34357580e-04, -4.37358211e-02],
[ -1.06433917e+00,  2.12501638e+00,  3.05234239e+00, ...,
  3.86340190e-01, -1.48467347e-01, -5.69036842e-01],
[ -5.01145154e-01,  7.40255346e-01, -5.11839275e-01, ...,
  -4.18045489e-01,  6.86860180e-01, -8.96531475e-01]])
```

## Training Models

```
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import ExtraTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import accuracy_score

# create instances of all models
models = {
    'Logistic Regression': LogisticRegression(),
    'Naive Bayes': GaussianNB(),
    'Support Vector Machine': SVC(),
    'K-Nearest Neighbors': KNeighborsClassifier(),
    'Decision Tree': DecisionTreeClassifier(),
    'Random Forest': RandomForestClassifier(),
    'Bagging': BaggingClassifier(),
    'AdaBoost': AdaBoostClassifier(),
    'Gradient Boosting': GradientBoostingClassifier(),
    'Extra Trees': ExtraTreeClassifier(),
}

for name, md in models.items():
    md.fit(X_train,y_train)
    ypred = md.predict(X_test)

    print(f"{name} with accuracy : {accuracy_score(y_test,ypred)}")

    Logistic Regression with accuracy : 0.9636363636363636
    Naive Bayes with accuracy : 0.9954545454545455
    Support Vector Machine with accuracy : 0.9681818181818181
    K-Nearest Neighbors with accuracy : 0.9590909090909091
    Decision Tree with accuracy : 0.9818181818181818
    Random Forest with accuracy : 0.9931818181818182
    Bagging with accuracy : 0.9886363636363636
    AdaBoost with accuracy : 0.1409090909090909
    Gradient Boosting with accuracy : 0.9818181818181818
    Extra Trees with accuracy : 0.9340909090909091

rfc =RandomForestClassifier()
rfc.fit(X_train,y_train)
ypred = rfc.predict(X_test)
accuracy_score(y_test,ypred)

0.9931818181818182
```

## Predictive System

```
def recommendation(N,P,k,temperature,humidity,ph,rainfall):
    features = np.array([[N,P,k,temperature,humidity,ph,rainfall]])
    prediction = rfc.predict(features).reshape(1,-1)

    return prediction[0]
```

```
N=50
P=60
k=50
temperature=30.0
humidity=20
ph=100
rainfall=100

predict=recommendation(N,P,k,temperature,humidity,ph,rainfall)

crop_dict ={1: "Rice", 2: "Maize", 3: "Jute", 4: "Cotton", 5: "Coconut", 6: "Papaya", 7: "Orange",
            8: "Apple", 9: "Muskmelon", 10: "Watermelon", 11: "Grapes", 12: "Mango", 13: "Banana",
            14: "Pomegranate", 15: "Lentil", 16: "Blackgram", 17: "Mungbean", 18: "Mothbeans",
            19: "Pigeonpeas", 20: "Kidneybeans", 21: "Chickpea", 22: "Coffee"}

if predict[0] in crop_dict:
    crop = crop_dict[predict[0]]
    print("{} is a best crop to be cultivated ".format(crop))
else:
    print("Sorry we are not able to recommend a proper crop for this environment")

    Apple is a best crop to be cultivated

import pickle
pickle.dump(rfc,open('model.pkl','wb'))

crop['label'].value_
```